



Getamazednow AI

A Getamazednow Company

ENTERPRISE ARCHITECTURE · EXECUTIVE WHITEPAPER

Agentic Enterprise Architecture

*A Multi-Cadence Operating Model
for Continuous Governance*



Prepared for Architecture Boards, Risk & Commercial Committees, and Executive Leadership

Enterprise Software · Cloud · Integration · Data & AI

Version 1.2 · Status: **FOR PEER REVIEW**



Contents

Contents	2
Executive Summary	3
1. The Problem: Cadence-Bound Governance	3
2. The Shift: From Documents to a Running Practice	4
3. The Operating Model: Five Loops, One Mesh	4
3.1 Daily — Sense	4
3.2 Weekly — Review	4
3.3 Monthly — Adjust	5
3.4 Quarterly — Realign	5
3.5 Yearly — Reset	5
4. Artefact-to-Cadence Framework	5
5. Agent Architecture	6
6. Stakeholder Engagement Model	7
7. Governance, Risk and Audit Considerations	7
8. Business Value Case	8
9. Adoption Roadmap	8
10. Conclusion	9

Executive Summary

Enterprise Architecture has historically run on committee time — quarterly boards, annual blueprint refreshes, point-in-time landscape audits. Technology change no longer runs on that clock. This whitepaper proposes an operating model in which a mesh of specialised agents runs the EA practice continuously, at five distinct cadences — Daily, Weekly, Monthly, Quarterly and Yearly — with each cadence producing signal for the next and each slower cadence setting constraints for the ones beneath it.

Twelve core artefacts — from business capability maps to the enterprise architecture blueprint — are distributed across these five loops according to their natural rhythm of change. Thirteen agents, coordinated by a single orchestrator, generate, maintain and govern these artefacts. Solution architects, engineering managers, product owners, heads of function, general managers, CxO stakeholders, external consultants and auditors each engage the mesh at the cadence appropriate to their decisions.

The mesh does not replace governance judgement. It removes the latency between a change occurring and that change being seen, assessed and acted on.

The intent is not to remove humans from EA decision-making — quarterly and yearly loops remain explicitly human-decision points — but to ensure that by the time a decision reaches a human, it is grounded in continuously current data, applied consistently, and fully traceable.

1. The Problem: Cadence-Bound Governance

Most EA practices inherit their operating rhythm from the governance calendar rather than from the pace of technology change. Landscape diagrams are refreshed for an audit and are stale within weeks. Portfolio reviews happen annually while the underlying asset health changes monthly. Governance gates are applied by different reviewers with different interpretations of the same principle. By the time an exception reaches the architecture board, the context that produced it may no longer be current.

This is not a failure of the people running EA — it is a structural mismatch between a document-cadence practice and a continuously changing landscape. The result is common across regulated and unregulated environments alike:

- Landscape maps that reflect what was documented, not what is deployed
- Portfolio decisions (Tolerate / Invest / Migrate / Eliminate) made on out-of-date cost and health data
- Governance exceptions accumulating without a clear signal that a principle itself needs revision
- Provenance for a governance decision living in an individual's memory rather than an auditable record

2. The Shift: From Documents to a Running Practice

The core design shift is to stop treating EA artefacts as documents produced on a calendar, and start treating them as the observable state of a continuously running practice. In this model, agents are not used to generate better documents faster — they are used to run the practice itself, with documents as one of several outputs.

This requires distinguishing between five distinct clock speeds, each with its own decision rights:

- Daily — telemetry and automated governance gates; no human decision unless a threshold is breached
- Weekly — design-time review and landscape accuracy; architects and engineering managers engage directly
- Monthly — tactical rebalancing of investment and delivery; product owners and heads of approve recommendations
- Quarterly — formal governance; general managers and risk committees make board-level decisions
- Yearly — strategic reset; CxO stakeholders set direction from a full year of grounded evidence

No loop operates in isolation. A daily governance breach can trigger an out-of-cycle quarterly review. A yearly blueprint decision cascades down as a new daily policy-as-code rule. This is what distinguishes an operating model from a reporting cadence.

3. The Operating Model: Five Loops, One Mesh

Each loop below is described by its purpose, the artefacts it primarily owns, and the human checkpoint at which agent output becomes a human decision.

3.1 Daily — Sense

Purpose: continuous telemetry with no human decision required unless a threshold is breached.

- Asset Health Assessments — the Health Agent scores every asset Red / Amber / Green from CMDB, APM, cloud cost, vulnerability scanner and end-of-life registry data
- Architecture Governance — the Gate Agent runs policy-as-code checks on new designs, pull requests, infrastructure-as-code templates and cloud deployments

Output routes to a human only on breach (alert to the relevant solution architect or engineering manager) or critical degradation (alert to the accountable Head of function).

3.2 Weekly — Review

Purpose: design-time governance and landscape accuracy.

- Technology Landscape Maps — the Landscape Agent diffs discovered infrastructure and integration state against the last known map, flagging drift such as shadow IT or undocumented integrations

- Architecture Assessments — the Assessment Agent runs structured, scored reviews of in-flight solution designs against current principles and patterns

Solution architects and engineering managers receive a weekly digest of assessments requiring their input and drift requiring explanation.

3.3 Monthly — Adjust

Purpose: tactical rebalancing of investment and delivery.

- Asset Portfolios — the Portfolio Agent runs a TIME analysis (Tolerate, Invest, Migrate, Eliminate) using health scores, cost, and strategic fit
- Asset Roadmaps — the Roadmap Agent re-baselines delivery roadmaps against actual programme progress
- Architecture Patterns — the Pattern Agent mines delivered solutions for reusable patterns and retires patterns no longer fit for purpose

Product owners and heads of function receive a portfolio pack for approval of TIME recommendations and roadmap adjustments — this is a decision point, not a status update.

3.4 Quarterly — Realign

Purpose: formal governance and strategic recalibration.

- Business Capability Maps — the Capability Agent refreshes the full capability heat-map against strategic priority
- Architecture Principles — the Principles Agent flags where repeated exceptions signal that a principle, not the enforcement of it, needs revision
- Strategy Transformations — the Transformation Agent tracks multi-quarter programmes against milestones and benefit realisation

This is the formal Architecture Board and Risk Committee cadence. Agents prepare the pack; general managers and the committee retain the decision.

3.5 Yearly — Reset

Purpose: strategic reset, synthesising a full year of grounded signal into direction-setting artefacts.

- Enterprise Architecture Blueprint — the Blueprint Agent synthesises capability, landscape, portfolio and principles into the annual blueprint
- Strategy White Papers — the Narrative Agent drafts thought-leadership papers from the year's transformation data and a market scan
- Strategy Transformations — next year's transformation initiatives are proposed based on capability gaps and portfolio outcomes

This is the highest-stakes, lowest-frequency loop, resolved at CxO strategy offsite. Agents present evidence-based options with trade-offs; direction remains a human decision.

4. Artefact-to-Cadence Framework

The following matrix maps all twelve core EA artefacts to their primary and secondary cadence. A solid marker indicates the cadence at which the artefact is created or substantially revised; a hollow marker indicates a cadence at which it is referenced, validated, or lightly updated.

Artefact	Daily	Weekly	Monthly	Quarterly	Yearly
Business Capability Maps			○	●	○
Technology Landscape Maps		●			
Asset Portfolios			●	○	
Asset Roadmaps			○	●	
Asset Health Assessments	●	○			
Architecture Principles				○	●
Architecture Patterns			●		
Architecture Governance	●	○		●	
Architecture Assessments		●			
Strategy Transformations				●	●
Strategy White Papers					●
Enterprise Architecture Blueprint				○	●

● Primary cadence ○ Secondary touch

5. Agent Architecture

Thirteen agents run beneath the five loops, coordinated by a single orchestrator that sequences the loops, routes outputs to the right stakeholder, and manages escalation across cadences.

Agent	Function
Orchestrator Agent	Sequences the five loops, routes outputs to the right stakeholder, manages escalation between cadences.
Health Agent	Continuous asset telemetry — EOL/EOS, vulnerability, performance monitoring.
Gate Agent	Policy-as-code enforcement on new designs, PRs, and deployments; maintains the exception register.
Landscape Agent	Infrastructure and integration discovery; flags drift from the documented map.
Assessment Agent	Structured, scored design reviews against principles and patterns.
Portfolio Agent	TIME analysis — Tolerate, Invest, Migrate, Eliminate — using health, cost and strategic fit.
Roadmap Agent	Re-baselines delivery roadmaps against actual programme progress.
Pattern Agent	Mines delivered solutions for reusable patterns; retires patterns no longer fit for purpose.
Capability Agent	Maintains the capability heat-map against strategic priority.

Agent	Function
Principles Agent	Detects principle drift from repeated exceptions; drives formal refresh.
Transformation Agent	Tracks multi-quarter transformation programmes against milestones and benefits.
Blueprint Agent	Synthesises capability, landscape, portfolio and principles into the annual blueprint.
Narrative Agent	Drafts thought-leadership white papers from the year's transformation data.

6. Stakeholder Engagement Model

A single shared dashboard is the wrong design goal. Each stakeholder group needs a different cadence entry point into the mesh, matched to the decisions they are actually accountable for.

Stakeholder	Cadence	Interaction Mode
Solution Architects	Daily / Weekly	Chat-based access to the Gate and Assessment Agents — design rationale returned with cited policy and pattern reference.
Engineering Managers	Daily / Weekly	Health and drift digests; automatic escalation on critical asset degradation.
Product Owners	Monthly	Portfolio and roadmap pack — approve TIME recommendations rather than only viewing them.
Heads Of	Monthly / Quarterly	Portfolio health and capability heat-map for investment sign-off.
General Managers	Quarterly	Board pack and risk committee view, agent-prepared and human-approved.
CxO Stakeholders	Quarterly / Yearly	Blueprint, white papers and transformation options presented with pros, cons and full provenance.
External Consultants	On-demand	Scoped, read-only access to the governance register, assessments and principles, with a full audit trail.
Auditors	On-demand	Immutable log of every gate decision, exception and override, with rationale attached, not just the outcome.

Every agent output carries provenance: the data source, the policy version, and the prior decision it built on. A governance decision an auditor cannot trace back to its evidence has no standing in a risk committee — this is a design requirement, not an afterthought.

7. Governance, Risk and Audit Considerations

Running EA as a continuously operating mesh changes where risk sits. The considerations below should be addressed explicitly before adoption, not discovered after.

Reward

- Continuous coverage — landscape and health state reflect reality daily, not at the last audit
- Consistent governance — the same principles are applied identically to every design, every time

- Full provenance — every decision cites its data source, policy version and prior precedent
- Faster cycle time — design reviews and gate checks move from board-cycle to same-day

Risk

- Automation complacency — auto-approvals can be rubber-stamped rather than genuinely assessed
- Principle ossification — rules continue to be enforced after they have stopped fitting the business
- Alert fatigue — continuous telemetry can overwhelm humans without thresholds tuned per asset criticality
- False confidence — agents optimise for what is measurable and can miss strategic or political nuance

Mitigations: sample-audit a percentage of auto-approvals monthly; use the quarterly principles review to examine exception patterns, not just counts; tune escalation thresholds per asset tier; keep the yearly loop an explicitly human decision with agents presenting options rather than conclusions; route low-confidence agent assessments to a human by default.

8. Business Value Case

The value of this model should be quantified against each organisation's own portfolio baseline before it is committed to a business case. The categories below indicate where the value is directional, and how each should be measured — they are not claimed figures.

- Governance cycle time — measure the median time from design submission to governance decision
- Landscape currency — measure the age of landscape data at the point decisions are made
- Portfolio discipline — measure the percentage of the portfolio carrying a current TIME classification
- Audit readiness — measure time-to-evidence when an auditor requests a decision trail
- Strategic coherence — measure variance between blueprint assumptions and realised portfolio outcomes
- Capacity redirection — measure architect and product owner hours spent on manual governance reporting, before and after

These measures let a committee assess the model on its own portfolio's evidence rather than an industry-wide claim.

9. Adoption Roadmap

Build the mesh bottom-up. A yearly Blueprint Agent with no daily telemetry beneath it is a faster document generator, not an operating model. The value is in the mesh, not any single loop.

Phase 1 (0–6 months) — Sense

Daily and Weekly loops only: Health and Gate Agents against one critical asset domain, Landscape and Assessment Agents for one delivery stream. Lowest risk, fastest signal on drift and breach.

Phase 2 (6–12 months) — Adjust

Add the Monthly loop: Portfolio and Roadmap Agents, TIME cycle run end-to-end with product owners, still fully human-approved.

Phase 3 (12–18 months) — Realign

Add the Quarterly loop: formal board pack automation, Capability and Principles Agents, first agent-prepared pack brought to the Architecture Board for critique.

Phase 4 (18–24 months) — Reset

Add the Yearly loop: full Blueprint Agent synthesis and Narrative Agent-drafted transformation proposals, presented at the CxO strategy offset.

10. Conclusion

An agentic Enterprise Architecture practice is not a faster way to produce the same documents. It is a change in what EA governance is made of — moving from a set of artefacts refreshed on a calendar to a continuously running mesh whose state those artefacts merely describe. The five-loop structure keeps the decision rights exactly where they belong: automated at the daily gate, human at the architecture board, and strategic at the CxO offsite — with full provenance connecting all three.

Start with one loop. Prove the signal. Build the mesh.